

Final Exam: CS 35L Section 7

University of California, Los Angeles

Fall 2017

Name: _____

Student ID: _____

Question	Points
1	/12
2	/16
3	/16
4	/16
5	/20
6	/20
Total	

Instructions:

1. This examination contains 9 pages in total, including this page.
2. You have **three (3) hours** to complete the examination.
3. Write your answers in this exam paper. The draft papers will not be accepted.
4. You may use any printed resources, including lecture notes, books, slides, assignments. **You cannot use any electronics.**
5. Please finish the final **independently**. If you have any question, please raise your hand.
6. There will be partial credit for each question.
7. Best of luck!

Question 1: Concept Question

[12 pts] Choose **two (2)** of the following three questions and answer them concisely.

- (a) What is cloud computing? What is the advantage (list at least two) of it?
- (b) What are the different levels (distance to hardware) of programming languages? Why do we need each of them?
- (c) Draw a simple diagram of how public-key encryption system works (what are the keys involved, how to use keys for encryption/decryption).

Average/Median/Max: 10.67/12/12

We do well in this question.

This is a relatively easy question. Copy the concepts from the slides and we receive the full credit.

Question 2: Unix Command

[16 pts]

The following is the subset of the man page for the “df” command.

DF(1) BSD General Commands Manual DF(1)

NAME

df -- display free disk space

SYNOPSIS

```
df [-b | -h | -H | -k | -m | -g | -P] [-ailn] [-t] [-T type]
    [file | filesystem ...]
```

LEGACY SYNOPSIS

```
df [-b | -h | -H | -k | -m | -P] [-ailn] [-t type] [-T type] [file |
    filesystem ...]
```

DESCRIPTION

The df utility displays statistics about the amount of free disk space on the specified filesystem or on the filesystem of which file is a part. Values are displayed in 512-byte per block counts. If neither a file or a filesystem operand is specified, statistics for all mounted filesystems are displayed (subject to the -t option below).

The following options are available:

- a Show all mount points, including those that were mounted with the MNT_IGNORE flag.
- b Use (the default) 512-byte blocks. This is only useful as a way to override an BLOCKSIZE specification from the environment.
- g Use 1073741824-byte (1-Gbyte) blocks rather than the default. Note that this overrides the BLOCKSIZE specification from the environment.
- i Include statistics on the number of free inodes. This option is now the default to conform to Version 3 of the Single UNIX Specification (‘‘SUSv3’’) Use -P to suppress this output.
- k Use 1024-byte (1-Kbyte) blocks, rather than the default. Note that this overrides the BLOCKSIZE specification from the environment.
- l Only display information about locally-mounted filesystems.
- m Use 1048576-byte (1-Mbyte) blocks rather than the default.
- n Print out the previously obtained statistics from the filesystems. This option should be used if it is possible that one or more filesystems are in a state such that they will not be able to provide statistics without a long delay. When this

option is specified, `df` will not request new statistics from the filesystems, but will respond with the possibly stale statistics that were previously obtained.

- P Use (the default) 512-byte blocks. This is only useful as a way to override an `BLOCKSIZE` specification from the environment.
- T Only print out statistics for filesystems of the specified types. More than one type may be specified in a comma separated list. The list of filesystem types can be prefixed with ‘no’ to specify the filesystem types for which action should not be taken. For example, the `df` command:

```
df -T nonfs,mfs
```

lists all filesystems except those of type NFS and MFS. The `lsvfs(1)` command can be used to find out the types of filesystems that are available on the system.

Given the above instructions, you should specify the sequence of commands you use to finish the following:

- (a) Use `df` to check the free disk space on the computer. The results should only include the locally-mounted filesystems, and exclude the `autofs` type filesystems. The display should use 1-Gbytes blocks.
- (b) Use pipeline and find all the lines that include the keyword “100%” in the previous output. Save the extracted results as `df.txt` to your HOME directory.
- (c) Go to your HOME directory. For `df.txt`, grant `r/w/x` permission to the owner, `r/x` permission to the group, and `r` permission to anyone.
- (d) Change the filename to `df_new.txt`, create a new directory called `test` under HOME directory, and copy the file into this directory.

Average/Median/Max: 13.57/15/16

We do well in this question as well.

- (a) Use the following three options:
 - `-l` (for local filesystems)
 - `-T noautofs` (for excluding `autofs`)
 - `-g` (for 1-G blocks)
- (b) We should use pipeline (`|`), ‘`grep`’, and redirection (`>`) in the answer. The Home directory should either be `~` or `$HOME`.
- (c) The correct syntax is ‘`chmod number filename`’. Some of us miss the filename.
- (d) Note that the question asks us to copy the file, instead of moving it. The syntax for `cp` is ‘`cp filename location`’.

Question 3: Bash Programming

[16 pts] Write a bash script “easy” to finish the following tasks:

- (a) The script takes two command line arguments (i.e. one extra parameter except for your script name). The second parameter is the path of a directory.
- (b) This script will check whether the input parameter is in the \$PATH variable. If not, add this directory into the \$PATH variable.
- (c) Under this directory, find the oldest regular file and gets the length of its filename. Denote this length as x .
- (d) Under this directory, find the number of files that is two-character long. Denote this number as y .
- (e) Write a function which calculates $x^3 + y^2 + x \cdot y$. Call this function and print the result to the standard output.
- (f) Your script should handle two exceptions and outputs “error”: when the input parameter number does not satisfy the requirement, and when the second argument is not a valid directory.

Average/Median/Max: 10.53/11.75/15

The outcome of this question is not as good as I expected. All the answers could be found either in notes or in homework, but I guess that we need to practice writing Bash more.

- (a) In Bash, we use \$0, \$1, \$2 ... to get the command line arguments.
- (b) The correct way to check this is to get the \$PATH variable, replace the : with space, and loop this list and compare with the second parameter. This can be found in Notes 3. The way to write to \$PATH variable is using export, which can be found in the Notes 1. Some minor mistakes:
 - Some of us compare strings using eq, which is used for arithmetic comparison.
 - Some of us does not translate the :. Without this step \$PATH is not a list for traversal.
 - Some of us use ‘grep’ to find whether the parameter is in the \$PATH. Although this is a great idea, we can have false positives: e.g., \$PATH = /usr/bin while the second parameter is /usr.
- (c) There are multiple ways to get the oldest regular file. We can either use commands or use the Bash processing to get it. Then we shall use \${#string} to get the filename length. One common mistake is that some of us find the youngest file, instead of the oldest one.
- (d) There are multiple ways to solve this question. We can either use commands, or use Bash to traverse the files and find the files that are two-char long.
- (e) We are supposed to write a valid Bash function. Please refer to notes about the format of a Bash function, how to call a function, and how to get the return (see Notes 3). We do not do well in this question; many of us use the C-style approach to either write a function or call the function.
- (f) We can use \$# to get the number of arguments, and use -d to check whether the argument is a valid directory.
- (g) Some general mistakes:
 - We should use backticks for command calling.
 - We should use special syntax for arithmetic calculation.
- (h) Note that we still get partial credit if we made any mistake mentioned above.

Question 4: Python Programming

[16 pts] In this question, you shall write a Python script to calculate the Euler's totient function of an integer n . Do not be scared away from the question; It is easy after you read and understand the following explanation. Hope that you also learn something from this question.

In number theory, Euler's totient function counts the positive integers up to a given integer n that are relatively prime (or co-prime)¹ to n , which is denoted as $\varphi(n)$. As an example, $\varphi(9) = 6$ since 1, 2, 4, 5, 7, 8 are relatively prime to 9, while 3, 6, 9 are not. $\varphi(10) = 4$ since 1, 3, 7, 9 are relatively prime to 10, while 2, 4, 5, 6, 8, 10 are not.

There exists an efficient way to calculate $\varphi(n)$. The critical step is to calculate prime factorization of a given n , which means that we will express an integer as the multiplication of a set of prime integers. For example, $360 = 2 \times 2 \times 2 \times 3 \times 3 \times 5 = 2^3 \times 3^2 \times 5^1$. After we get this factorization, it is easy to calculate the $\varphi(n)$. Suppose that $n = p_1^{k_1} \cdots p_r^{k_r}$, where $p_1 \cdots p_r$ are different prime numbers, $\varphi(n) = n(1 - 1/p_1)(1 - 1/p_2) \cdots (1 - 1/p_r)$. For example, $10 = 2^1 \times 5^1$, and $\varphi(10) = 10 \times (1 - 1/2) \times (1 - 1/5) = 4$; $360 = 2^3 \times 3^2 \times 5^1$, and $\varphi(360) = 360 \times (1 - 1/2) \times (1 - 1/3) \times (1 - 1/5) = 96$.

Luckily enough, you are not asked to calculate the prime factorization of n . Instead, there exists a library called `zhaowei` that helps you do it. In this library, you can call a function called "`p-factor`", which takes an integer as the input, and outputs a dictionary which represents the prime factorization of the input.

Calls `p-factor(10)` returns a dictionary `{2:1, 5:1}`, since $10 = 2^1 \times 5^1$.
Calls `p-factor(360)` returns a dictionary `{2:3, 3:2, 5:1}`, since $360 = 2^3 \times 3^2 \times 5^1$.

Note that if the input for `p-factor()` is not an integer greater than 1, the output is an empty dictionary.

- (a) Describe how to install the package `zhaowei` using Python package management tools.
- (b) Write a Python 2 scripts that takes one parameter n from the standard input. Calculate and print the Euler's totient function $\varphi(n)$ of n . You shall utilize the library `zhaowei` and the function `p-factor`.
- (c) Your script should output "error" if the input is not a positive integer, and $\varphi(n)$ otherwise. You **do not** have to consider any other exceptions. Hint: $\varphi(1) = 1$ should be considered separately because `p-factor` only works for integer greater than 1.

Average/Median/Max: 11.27/12.25/15.5

The outcome of this question is good. If you read through the description carefully, it is a short program within 20 lines of codes.

- (a) Note that we are talking about Python package management here. The answer should either be `pip` or `easy_install`.
- (b) We should pay attention to the following things
 - We shall import the library `zhaowei`.
 - We should correctly read the parameter from command line, e.g. using `sys.argv`. Please refer to our Python homework.
 - We can iterate the returned dictionary by using something like `'for i in dict.keys()'`.
 - The calculation should be correct. Note that we are using the dictionary keys to do the calculation, not the values.

¹Two integers x and y are said to be co-prime if the only positive integer that divides both of them is 1. E.g., 10 and 9 are co-prime, while 10 and 8 are not. As a side note, obviously, any two prime numbers are co-prime.

- (c) This is a tricky question. The question asks us to print ‘error’ whenever the input is not a positive integer. However, most of us only consider when the input is is a non-positive integer. We should also consider the case where the input is non-integer rational numbers, or even irrational numbers (although they are not inputtable from standard input)! You can implement this by either directly judging whether the input is an integer, or you can utilize this feature of the function `p-factor()`:

“if the input for `p-factor()` is not an integer greater than 1, the output is an empty dictionary.”

In other words, we can put an if statement after we get the return from `p-factor()`: whenever the returned dictionary is empty, we should output “error”. Do not forget to handle the case where the input is 1 though.

- (d) Note that we still get partial credit if we made any mistake mentioned above.

Question 5: C Programming: Multithreading, Git

[20 pts] There exists a Git directory in the remote Git server. <https://github.com/ZhaoweiTan/final.git>. This Git project only has one branch **master**. Inside this directory is a single file called **vector.c**. Inside the file it reads:

```
double dot_product(double v[], double u[], int n) {
    int i;
    double result = 0;
    for (i = 0; i < n; i+=1)
    {
        result += v[i]*u[i];
    }
    return result;
}
```

You are asked to do the following. Your answer should include the commands you use and the program you write.

- Git clone the code to your local computer under \$HOME/test directory. (Suppose the directory is already there, but you have to change directory first)
- Create a new branch called “test” and switch to it.
- In this new branch, create a new file called **mt_vector.c**.
- Understand the code in **vector.c**; write a multithreading version of it in **mt_vector.c**. There should be at least a function called **dot_product_multithread** in your program, which utilizes **four (4)** threads and returns the identical results. When an error occurs when handling pthread, prints “error” and returns 999.
- Commit the changes and merge the changes to the **master** branch.

Average/Median/Max: 14.53/16.25/20

The outcome of this question is good. We are familiar with C programming, but we have to be more careful when writing the code.

Some general comments on git:

- All the commands that we are using can be found in the Slides. Note that for commit changes we should first ‘add’ then ‘commit’. For merging, we should first switch back to the master branch and then merge the changes.

Some general comments on multithreading:

- We need to use the correct pthread function for creating and joining. Remember to check whether these functions return error whenever we use them.
- The parameter passing is crucial. We have to make sure the parameter is passed correctly from the main function to multiple threads. The safest way is to create multiple global variables, each of which store the subset of the initial vector that will later be processed by a thread.
- We need to add up the calculated value from multiple threads and collectively return the correct result.

Question 6: C Programming: System Call and Compiling

[20 pts]

- (a) Firstly, please list the four compiling stages, and explain what each step roughly does.
- (b) Suppose that you have a file called `input.txt` under your HOME directory. Write a `main.c` script. In the `main` function use system calls to read the content in this file.
- (c) Suppose that the content in `input.txt` is a string. Write a C file called `reverse.c` and its header `reverse.h`. In these files write a function `reverse_string` that reverses a string.
- (d) In `main.c`, call the function in `reverse.c` on the string you just read from `input.txt`. You should dynamically load `reverse_string` function in `reverse.c` using `dlopen` and `dlsym`. You are supposed to close up in the end as well.
- (e) Use system call to write the reversed string to the standard output.
- (f) Your program should consider all the linking errors and system call errors. If an error occurs, simply returns -1. Otherwise, a successful `main` should return 0.
- (g) Write a simple makefile that compiles your library and program.
- (h) What command can check the dynamic libraries your program uses?

Average/Median/Max: 15.28/16.75/20

The outcome of this question is good. We are familiar with C programming, but we have to be more careful when writing the code.

Some general comments on multithreading:

- We have to use the following system calls: `open`, `read`, `write` and `close` to open the file, read from that file, write to standard output, and close the file respectively. After each system call, we should check the return value in case there is error.
- For dynamic loading we should use `dlopen` and `dlsym`. To close up we use `dlclose`. Remember we need to check error after each of this.
- The Makefile should be working. Write something similar to Assignment 8 would be good enough.